



Pierre Pfister | VPP Infrastructure Libraries

FD.io /dev/boot Training

31 May - 03 June | Télécom ParisTech



VPP Infrastructure Libraries

Warning: Extreme C coding



All the code, structure and libraries

VPP Programming FOR DUMMIES®

Forward MORE
MPPs !!!

*A Reference
for the
Rest of Us!*

By FD.io



vppinfra structures



Sub-agenda

types.h

We use:

u8, u16, u32, u64, u32x4 (intrinsics)

i8, ...

f32, f64


uword

vec.h

A vector is an auto-resized C array with some meta data:

Parameters: Type T, header size, element alignment

```
T e[5];  
//Default alignment  
//No header  
T *v1 = 0; //Init  
vec_add1(v1, &e[0]); //Add 1  
vec_add2(v1, &e[1], 4); //Add 4  
v1
```



```
//Header of 8 bytes  
//Alignment of 16 bytes  
T *v2 = 0; //Init  
vec_add2_ha(v2, &e[0], 5, 8, 16);  
vec_header(v2, 8);
```



```
vec_len(v1); //length safe when v1==0  
_vec_len(v1) = 0; //length unsafe
```

```
//increase length to lth elmt.  
//Set new elements to 0.  
vec_validate(V,I);
```

```
// Allocate for N more elements (but does not change length)  
vec_alloc(V,N);
```

```
vec_free(v); //Free memory and set v to 0  
vec_reset_length(v); //set length to 0 (safe)
```

Allocation only increases.

Vector origin pointer may change !

Store Indexes (not pointers) !



bitmap.h

Set and get bits with indexes. Lots of them.

Implemented as a uword vector.

```
uword *bitmap = 0;

//Allocate bitmap for 100 bits
clib_bitmap_alloc(bitmap, 100);

//Make room for 201 bits
clib_bitmap_validate(bitmap, 200);

clib_bitmap_set(bitmap, 33, 1);
clib_bitmap_get(bitmap, 33);

//Get bits from 3 to 12 included
uword a = clib_bitmap_get_multiple
          (bitmap, 3, 10);

clib_bitmap_set_multiple ...
```

vec	length
0..63	b
64..127	b
...	

```
uword clib_bitmap_first_set (uword * ai);
uword clib_bitmap_first_clear (uword * ai);

uword clib_bitmap_next_set (uword * ai, uword i);
uword clib_bitmap_next_clear (uword * ai, uword i);

uword clib_bitmap_count_set_bits (uword * ai);

...
```

pool.h

Fixed sized element allocator.

Based on a vec and a bitmap.

```
T *pool = 0;
```

```
T *e;
```

```
pool_get(pool, e);
```

```
//pool_get_aligned(pool, e, 4);
```

```
//Query if element is free
```

```
pool_is_free(pool, e);
```

```
pool_is_free_index(pool, e - pool);
```

```
//Unalloc element
```

```
pool_put(pool, e);
```


```
//Free the whole pool
```

```
pool_free(pool);
```

```
// Allocated element count
```

```
pool_elts(pool);
```

```
typedef struct {  
    uword * free_bitmap;  
    u32 * free_indices;  
} pool_header_t;
```



Free bitmap

Keeps track of element allocation

Free indices

Used for fast re-allocation

Sparse memory assignment.

Avoids memory fragmentation.

Fast.

heap.h

Variable element size allocator.

```
T *heap = 0;
T *e;

//Allocate 4xT.
U32 handle;
u32 offset = heap_alloc(heap, 4, handle);
//Allocated heap[offset] - heap[offset + 4]

//Object size
heap_size(heap, handle);

//Deallocate
heap_dealloc(heap, handle);
```

Rarely used (pools are faster).

Still efficient.

To be used if you need variably sized allocations.

e.g. classifier

bihash.h

Hashing table on steroid.

Heart of the ipv6 fib

bihash_8_8.h Key size
 Value size
bihash_24_8.h

bihash_template.h Code with macro magic to add your own bihash
(define BIHASH_TYPE as X_Y and include the header)

vppinfra

Parse'n'print



Sub-agenda

format.h

Printing and Scanning has never been so easy.

Format == write from data to strings

```
struct {  
    u32 v1, v2;  
} example;
```

```
u8 *format_example(u8 *s, va_list *args) {  
    struct example *e = va_arg(*va, struct example *);  
    return format(s, "(%d, %d)", e->v1, e->v2);  
}
```

```
u8 *format_example_with_u32(u8 *s, va_list *args) {  
    struct example *e = va_arg(*va, struct example *);  
    u32 v3 = va_arg(*va, u32);  
    return format(s, "%U with %d", format_example, e, v3);  
}
```

```
struct example e = {1, 2};  
u8 *str = format(0, "%U", format_example_with_u32, &e, 5);  
clib_warning("%s", str);  
vec_free(str);
```

```
typedef u8 * (format_function_t) (u8 * s, va_list * args);
```

```
u8 * va_format(u8 * s, char * format, va_list * args);  
u8 * format(u8 * s, char * format, ...);
```

Supports 'most' of printf syntax.



format.h

Printing and Scanning has never been so easy.

Unformat == get from string to data

```
typedef uword (unformat_function_t) (unformat_input_t *  
input, va_list * args);
```

```
struct {  
    u32 v1, v2;  
} example;
```

```
uword unformat_example (unformat_input_t *i, va_list *args) {  
    struct example *e = va_arg (*va, struct example *);  
    return unformat(i, "(%d, %d)", &e->v1, &e->v2);  
}
```

```
uword unformat_example_with_u32(unformat_input_t *i, va_list *args) {  
    struct example *e = va_arg (*va, struct example *);  
    u32 *v3 = va_arg (*va, u32 *);  
    return unformat(i, "%U with %d", unformat_example, e, v3);  
}
```

```
struct example e; u32 v;  
const char *str = "(1, 2) with 3 foo bar";  
unformat_input_t i;  
unformat_init_string(&i, str, strlen(str));  
unformat(i, "%U", unformat_example_with_u32, &e, &v);
```

```
uword unformat (unformat_input_t *i, char * fmt, ...);  
uword va_unformat (unformat_input_t *i, char * fmt,  
va_list * args);
```

Supports 'most' of scanf syntax.
Returns whether unformat was successful.

In case of success, the data is consumed
in the unformat_input.
Otherwise, the data is left untouched.



format.h

Printing and Scanning has never been so easy.

Parsing CLI input example

From vnet/vnet/l2tp/l2tp.c

Arguments can be in any order.

client <ip6_addr> our <ip6_addr> local-
cookie <long> remote-cookie <long> ...

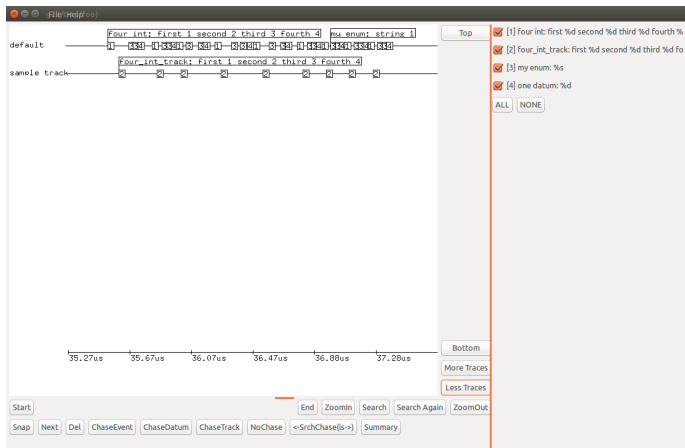
```
if (!unformat_user(input, unformat_line_input, line_input))
    return 0;

while (unformat_check_input(line_input) != UNFORMAT_END_OF_INPUT) {
    if (unformat(line_input, "client %U",
                 unformat_ip6_address, &client_address))
        client_address_set = 1;
    else if (unformat(line_input, "our %U",
                     unformat_ip6_address, &our_address))
        our_address_set = 1;
    else if (unformat(line_input, "local-cookie %llx", &local_cookie))
        ;
    else if (unformat(line_input, "remote-cookie %llx", &remote_cookie))
        ;
    else if (unformat(line_input, "local-session-id %d",
                     &local_session_id))
        ;
    else if (unformat(line_input, "remote-session-id %d",
                     &remote_session_id))
        ;
    else if (unformat(line_input, "l2-sublayer-present"))
        l2_sublayer_present = 1;
    else
        return clib_error_return(0, "parse error: '%U'",
                                format_unformat_error, line_input);
}
```



vppinfra

Event Logger

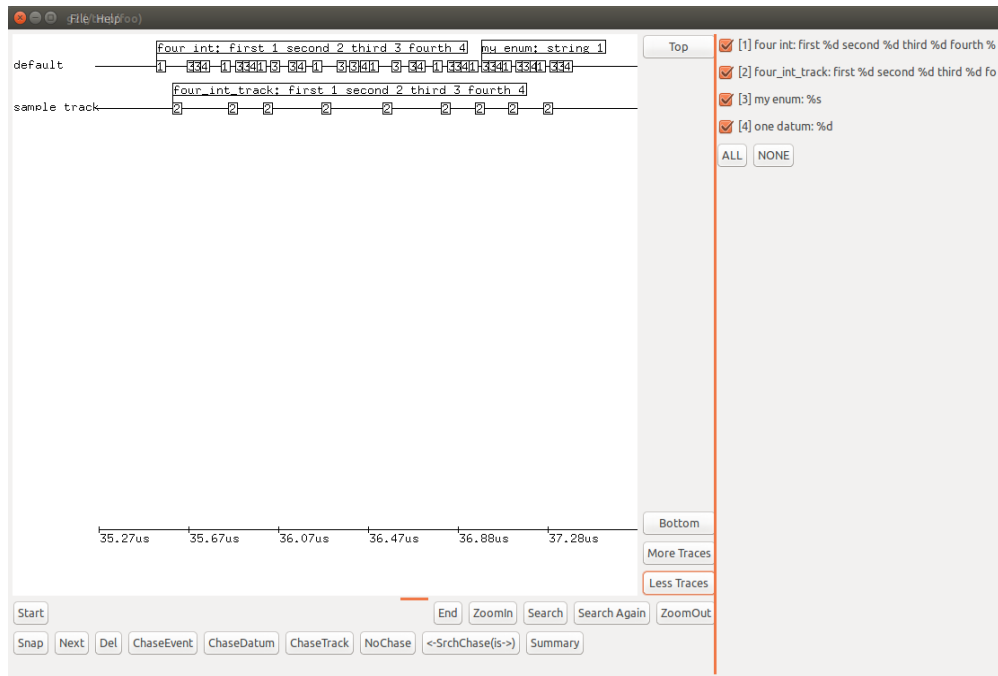


Sub-agenda

eelog.h

High speed event logging

<https://wiki.fd.io/view/VPP/eelog>



- Event-logging enabled in `.../vlib/vlib/main.c:vlib_main(...)`
- Use the `eelog_main_t` in `vlib_global_main`, aka `&vlib_global_main.eelog_main`
- Default ring size 128K events
- Thread safe—lock-free atomic increment to dole out event slots
- Each event-slot is 32 bytes: u64 time-in-cpu-clocks, u16 event-id, u16 track, 20 bytes of data
- Logging an event costs less than 100ns
- Observer effect: at most a couple of events per node, per frame at speed

Cf. Dave Barach's presentation:

https://docs.google.com/presentation/d/1C_1zM5Z3sTibOj1e2pe_YDbiMCytwZspK541fPElyWM



vppinfra
The rest




Sub-agenda

Things you may or may not learn by yourself

fheap.h	Fibonacci heap...
graph.h	Graph Implementation
hash.h	Hash table
phash.h	Again
qhash.h	Again
longjmp.h	We actually use it !

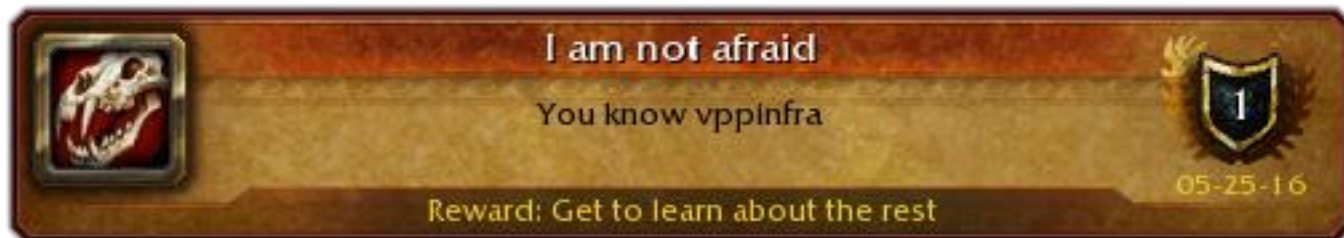
Etc...



Lessons from vppinfra

1. Store indexes, not pointers (arrays get resized).
2. Macros do change parameters (and not only pointers).
3. Format and unformat functions are usefull.
4. Aligement can be important for multi-thread efficiency.

Congratulation !



vlib



Sub-agenda

Vlib Overview

VPP is based on 2 main ideas:

1. Vectors (0 to 256 buffers bundled together)
2. Nodes (in a graph) performing actions and passing the vectors.

vlib is where all this abstract model is defined. Outside of any networking considerations (No IP, no ethernet, etc...).

It also includes:

- Counters
- CLI
- Scheduler (main loop).
- A few other things...



Vlib nodes

i-am-a-vlib-node

Sub-agenda

Vlib nodes - node.h – node_funcs.h

Nodes have:

- A name dpdk-rx, ip6-forward, tapcli-tx, ...
- An index Uniquely identifies a node
- A function Node callback to operate on vectors
- A type
- A set of next nodes Identified by index & counters
- A set of errors With names & counters

Each node has a registration data (`vlib_node_registration_t`) and a runtime data (`vlib_node_t`).

Nodes are created with constructors at initialization.

It is possible to dynamically add next nodes.

always_inline uword

vlib_node_add_next (**vlib_main_t** * vm, **uword node,**
uword next_node);

```
typedef enum {  
    VLIB_NODE_TYPE_INTERNAL,  
    VLIB_NODE_TYPE_INPUT,  
    VLIB_NODE_TYPE_PRE_INPUT,  
    VLIB_NODE_TYPE_PROCESS,  
    VLIB_N_NODE_TYPE,  
} vlib_node_type_t;
```

```
VLIB_REGISTER_NODE (l2input_node) = {  
    .function = l2input_node_fn,  
    .name = "l2-input",  
    .vector_size = sizeof (u32),  
    .format_trace = format_l2input_trace,  
    .format_buffer = format_ethernet_header_with_length,  
    .type = VLIB_NODE_TYPE_INTERNAL,  
    .n_errors = ARRAY_LEN(l2input_error_strings),  
    .error_strings = l2input_error_strings,  
    .n_next_nodes = L2INPUT_N_NEXT,  
    .next_nodes = {  
        [L2INPUT_NEXT_LEARN] = "l2-learn",  
        [L2INPUT_NEXT_FWD] = "l2-fwd",  
        [L2INPUT_NEXT_DROP] = "error-drop",  
    },  
};
```



Vlib nodes - *VLIB_NODE_TYPE_INTERNAL*

Most typical node receiving buffer vectors, performing operations.

Includes tx nodes.

Vlib nodes - *VLIB_NODE_TYPE_INPUT*

Typically device input nodes.

Create frames from scratch and dispatch to internal nodes.

```
typedef enum {  
    VLIB_NODE_STATE_POLLING,  
    VLIB_NODE_STATE_INTERRUPT,  
    VLIB_NODE_STATE_DISABLED,  
    VLIB_N_NODE_STATE,  
} vlib_node_state_t;
```

Constantly called

Called once when interrupted

Not called

u32

input_main_loops_per_call;

Throttle polling

Vlib nodes - *VLIB_NODE_TYPE_PRE_INPUT*

Called before input nodes.

Not used as far as I know.



Vlib nodes - *VLIB_NODE_TYPE_PROCESS*

Thread-like function.

Can be suspended, wait for events, be resumed...
(based on setjump/longjump).

Wait for an event

always_inline f64

vlib_process_wait_for_event_or_clock (
 vlib_main_t * vm, f64 dt)

Send an event

always_inline void

vlib_process_signal_event (vlib_main_t * vm,
 uword node_index,
 uword type_opaque, uword data);

```
static uword ip6_icmp_neighbor_discovery_event_process(vlib_main_t * vm,  
vlib_node_runtime_t * node, vlib_frame_t * frame)  
{  
    uword event_type;  
    ip6_icmp_neighbor_discovery_event_data_t * event_data;  
    while (1) {  
        vlib_process_wait_for_event_or_clock(vm, 1. /* seconds */);  
        event_data = vlib_process_get_event_data(vm, &event_type);  
        if(!event_data) {  
            ip6_neighbor_process_timer_event(vm, node, frame);  
        } else {  
            switch(event_type) {  
                case ICMP6_ND_EVENT_INIT:  
                    break;  
                case ~0:  
                    break;  
                default:  
                    ASSERT(0);  
            }  
            if(event_data)  
                _vec_len(event_data) = 0;  
        }  
    }  
    return frame->n_vectors;  
}
```

Vlib/unix



NO IMAGE
AVAILABLE

Sub-agenda

Vlib and file descriptors

Vlib implements a poll – event-loop-like node.
Listens for fd events and execute callbacks.
Somehow slow due to `node->input_main_loops_per_call`

```
typedef struct unix_file {  
    u32 file_descriptor;  
  
    u32 flags;  
#define UNIX_FILE_DATA_AVAILABLE_TO_WRITE (1 << 0)  
#define UNIX_FILE_EVENT_EDGE_TRIGGERED (1 << 1)  
  
    uword private_data;  
  
    unix_file_function_t * read_function, * write_function, *  
    error_function;  
} unix_file_t;
```

```
uword unix_file_add (unix_main_t * um, unix_file_t * template);
```

Example from `vnet/unix/tapcli.c`

```
unix_file_t template = {0};  
template.read_function = tapcli_read_ready;  
template.file_descriptor = dev_net_tun_fd;  
ti->unix_file_index = unix_file_add (&unix_main,  
&template);
```

```
static clib_error_t * tapcli_read_ready (unix_file_t * uf)  
{  
    vlib_main_t * vm = vlib_get_main();  
    tapcli_main_t * tm = &tapcli_main;  
    uword * p;  
    vlib_node_set_interrupt_pending (vm, tapcli_rx_node.index);  
    ...  
    return 0;  
}
```



See more at [vlib/unix/unix.h](https://vlib.unix.unix.h)

Plugins vlib/unix/plugin.h

Plugins are loaded at VPP init:

```
$ vpp plugin_path <path>
```

Vpp tries to load all libraries present in the path.

`vlib_plugin_register` must exist and return 0.

'constructor' functions are also executed.

```
clib_error_t * constr_example (vlib_main_t * vm);  
VLIB_INIT_FUNCTION (constr_example);
```

static int

```
load_one_plugin (plugin_main_t *pm, plugin_info_t *pi, int from_early_init)  
{  
    void *handle, *register_handle;  
    clib_error_t * (*fp)(vlib_main_t *, void *, int);  
    clib_error_t * error;  
    void *handoff_structure;  
  
    handle = dlopen ((char *)pi->name, RTLD_LAZY);  
    [...]  
    pi->handle = handle;  
  
    register_handle = dlsym (pi->handle, "vlib_plugin_register");  
    if (register_handle == 0) {  
        dlclose (handle);  
        return 0;  
    }  
    fp = register_handle;  
    handoff_structure = vnet_get_handoff_structure();  
    if (handoff_structure == 0)  
        error = clib_error_return (0, "handoff structure callback returned 0");  
    else  
        error = (*fp)(pm->vlib_main, handoff_structure, from_early_init);  
  
    if (error) {  
        clib_error_report (error);  
        dlclose (handle);  
        return 1;  
    }  
  
    clib_warning ("Loaded plugin: %s", pi->name);  
    return 0;  
}
```



CLI

To be explained in more details in CLI and API session

Vlib buffers



Or not 

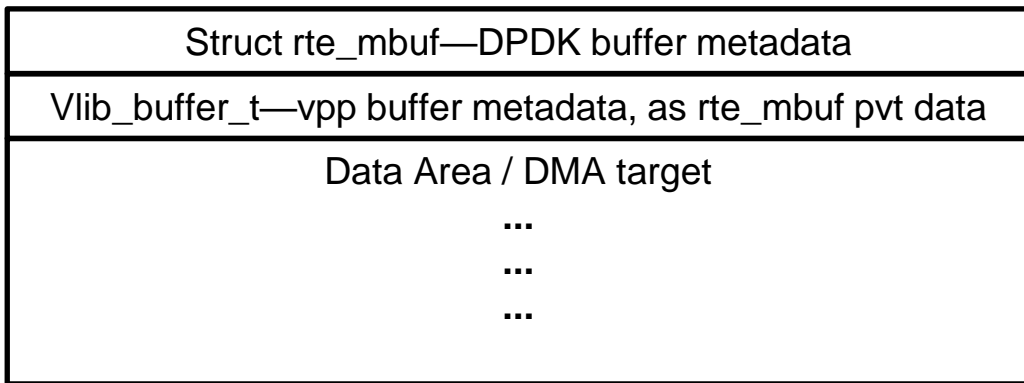
That is the question

Sub-agenda

Vlib buffers – vlib/buffer.[ch] vlib/dpdk_buffer.c

`vlib_buffer_t` is VPP's 'packet' basic structure (equivalent to DPDK's `rte_mbuf`).

When `DPDK == 1`, `vlib_buffer_t` is encapsulated within `rte_mbuf`.



```
#define vlib_buffer_from_rte_mbuf(x) ((vlib_buffer_t *) (x+1))
```

```
#define rte_mbuf_from_vlib_buffer(x) (((struct rte_mbuf *) x) - 1)
```

Vlib buffers

i16 current_data;

u16 current_length;

u32 flags;

Signed offset in data[], pre_data[]

Nbytes between current data and the end of this buffer.

```
#define VLIB_BUFFER_IS_TRACED (1 << 0)
#define VLIB_BUFFER_LOG2_NEXT_PRESENT (1)
#define VLIB_BUFFER_NEXT_PRESENT \
    (1 << VLIB_BUFFER_LOG2_NEXT_PRESENT)
#define VLIB_BUFFER_IS_RECYCLED (1 << 2)
#define VLIB_BUFFER_TOTAL_LENGTH_VALID \
    (1 << 3)
#define VLIB_BUFFER_HGSHM_USER_INDEX_VALID
    (1 << 4)
#define VLIB_BUFFER_REPL_FAIL (1 << 5)
#define LOG2_VLIB_BUFFER_FLAG_USER(n)
    (32 - (n))
```

Vlib buffers

u32 free_list_index;

Barely used

u32 total_length_not_including_first_buffer;

End of vlib_buffer_init_for_free_list() init

u32 next_buffer;

Buffer index of next buffer (if present)

u32 trace_index;

Valid if buffer traced

u32 clone_count;

If non-zero, xmit copy and recycle the original buffer
u16 error_node, error_code – set to arrange
counter bump

vlib_error_t error;

u32 opaque[8];

subgraph metadata, see .../vnet/vnet/buffer.h

Vlib buffers

u32 opaque2[16];

Barely used

u8 pre_data[VLIB_BUFFER_PRE_DATA_SIZE];

Rewrite space, 128 bytes

u32 data[0];

Aligned DMA target. Certain hardware devices DMA into data[n], e.g. n=6.

Vlib buffers

`vnet_buffer_opaque_t` – union of types which contain ‘subgraph’ metadata

- 8x32 bits
- `vnet_buffer(b)->sw_if_index[VLIB_RX, VLIB_TX]` are common
- From driver level: `vnet_buffer(b)->sw_if_index[VLIB_RX]` set to RX (physical) interface
`vnet_buffer(b)->sw_if_index[VLIB_TX] = ~0`
- ip4/6-lookup use `VLIB_TX` as FIB index override.
- These are like node input/output metadata (e.g. ip-frag).



Vnet_buffer_opaque_t intersection analysis

You nailed it !

